

[Date]

AI's 10 to Watch Selection Committee
IEEE Intelligent Systems

Dear Members of the Selection Committee,

I am pleased to strongly recommend Rahul Gupta for IEEE Intelligent Systems' AI's 10 to Watch. I have worked with Rahul for almost six years on research at the intersection of trustworthy natural language processing, responsible AI, fairness, red-teaming, and foundation-model evaluation. During that time, I have seen him develop into one of the rare industry scientists who can both generate novel research ideas and build the organizational structures needed to make those ideas useful to the broader AI community.

My recommendation is based on direct collaboration. Rahul and I have coauthored and collaborated on a sustained body of work in responsible NLP and trustworthy AI, including BOLD: Dataset and Metrics for Measuring Biases in Open-Ended Language Generation; Measuring Fairness of Text Classifiers via Prediction Sensitivity; Incorporating Fairness in Large Scale NLU Systems; An Analysis of the Effects of Decoding Algorithms on Fairness in Open-Ended Language Generation; Resolving Ambiguities in Text-to-Image Generative Models; FLIRT: Feedback Loop In-context Red Teaming; The Steerability of Large Language Models toward Data-Driven Personas; and recent work on detoxification, data curation, and model unlearning. We are also co-inventors on Synthetic Data Generation for Machine Learning Models. This long collaboration gives me a concrete basis for judging Rahul's technical creativity, research judgment, and impact.

Rahul's central strength is his ability to balance industrial innovation with genuine support for academic research. Many industry research leaders are effective at translating academic ideas into products, and many academics are effective at advancing open scientific questions. Rahul stands out because he works seriously in both directions. He identifies practical problems from deployed AI systems, turns them into research questions that are intellectually meaningful, and then creates mechanisms for academics and practitioners to study those problems rigorously.

I have repeatedly found Rahul to be ingenious in how he formulates new research directions. In our joint work on BOLD, he helped connect fairness measurement to the realities of open-ended generation, where harms cannot be reduced to a single classification metric. In our work on fairness in text classifiers and large-scale NLU systems, he brought a clear view of how responsible AI methods must survive the complexity of real production systems. In later work on FLIRT, data curation, steerability, and unlearning, he again showed a talent for turning emerging concerns in foundation models into concrete evaluation and mitigation problems that researchers can work on. [Kai-Wei can add a brief specific anecdote here about a research idea Rahul initiated or reframed.]

Equally important, Rahul has translated these research ideas into industrial mechanisms. At Amazon, he founded and led Responsible AI science efforts for Alexa, later led Responsible AI for Amazon Nova, and now serves as the overall Responsible AI lead for Amazon's AGI / Frontier AI Assets organization. His patented and patent-pending technologies provide concrete evidence of this translation. In Alexa, his patent portfolio includes inventions around cohort determination, model configuration, user-data processing, spoken-language understanding, and NLP component evaluation. These technologies helped make Alexa experiences more equitable, fair, and private for large populations of users. In the Nova and foundation-model context, Rahul's patented work and filings around generative language models, synthetic-data generation, and responsible AI methods have supported improvements in model evaluation and safety profile.

What makes Rahul especially valuable to the AI community is that he does not treat industry impact and academic impact as competing goals. He has built platforms that help other researchers gain access, visibility, support, and recognition. I have seen this directly through his organization of the

TrustNLP workshop series and related community efforts in trustworthy NLP. Rahul has also helped organize challenge-style research programs, including the SemEval Unlearning Challenge, that give academic researchers shared tasks, benchmarks, and public venues for comparing ideas. Through TrustNLP, Trustworthy Speech Processing, the SemEval Unlearning Challenge, and the Nova Trusted AI Challenge, Rahul has helped create infrastructure where researchers can test ideas, publish results, and receive community recognition. The Nova Trusted AI Challenge is a particularly strong example: it attracted more than 100 university applications and supported 10 university teams through multi-million-dollar awards for frontier-model safety research. [Kai-Wei can add a sentence here if he directly observed Rahul's role in shaping one of these academic programs.]

This kind of contribution is especially important for AI's 10 to Watch. Rahul is not following a conventional academic path, where recognition is measured primarily by faculty awards or individual grants. As an industry scientist, he has had fewer opportunities to apply for many academic early-career awards. Instead, his recognition has come through increasing scientific responsibility, internal research investment, patents, deployed systems, and research programs that support the academic community. In my view, that should be seen as a strength rather than a gap. Rahul has built the kind of field-shaping bridge between industry and academia that the AI community increasingly needs.

Across the years I have known him, Rahul has consistently demonstrated unusual technical judgment, creativity, and generosity toward collaborators. He is able to see around corners: to recognize which responsible AI problems will matter before they become standard benchmarks, to identify evaluation gaps in emerging systems, and to mobilize teams and collaborators around tractable research formulations. He also has the temperament of a community builder. He supports academic practitioners, creates opportunities for students and faculty, and treats responsible AI not simply as an internal product requirement but as a shared scientific challenge.

For these reasons, I believe Rahul Gupta is an excellent candidate for AI's 10 to Watch. He has already made significant contributions to trustworthy AI research and practice, and his trajectory suggests even greater future impact as responsible AI becomes central to the development and deployment of frontier AI systems. I recommend him with enthusiasm.

Sincerely,

Kai-Wei Chang

Professor of Computer Science, UCLA

Amazon Scholar, Amazon AGI

Email placeholder