

Rahul Gupta

Responsible AI, Privacy, and Data Governance Leader

(213)587-1103
rahul.1987iit@gmail.com



EXECUTIVE PROFILE

- Enterprise AI governance leader in Amazon's AGI organization, accountable for Responsible AI programs spanning safety, privacy, fairness, red-teaming, evaluation, and mitigations for frontier models.
- Built and led cross-functional research and governance teams of 50+ people, partnering with science, product, engineering, legal, policy, academic, and executive stakeholders to turn emerging AI risks into practical operating mechanisms.
- Hands-on scientist with first-author work on [CBRN uplift evaluation for frontier language models](#) and [fairness in large-scale NLU systems](#); 25+ patented or patent-pending technologies and publications across NeurIPS, ICLR, ACL, EMNLP, FAccT, and ICASSP.

GOVERNANCE & TECHNICAL LEADERSHIP EXPERIENCE

Senior Applied Science Manager June 2026 – Present
Amazon / Frontier AI Assets, Responsible AI Cambridge, MA

- Overall head of Responsible AI in Amazon's newly architected AGI division, reporting directly to a Vice President.
- Set enterprise priorities for responsible frontier-model development, including model safety, privacy, fairness, evaluation, red teaming, policy alignment, and launch readiness.
- Translate ambiguous AI governance and risk questions into measurable frameworks, executive guidance, and scalable mechanisms for product and technology teams.

Senior Applied Science Manager October 2022 – May 2026
Amazon AGI / Nova Responsible AI Cambridge, MA

- Led Nova Responsible AI programs covering safety evaluations, red teaming, fairness, privacy-sensitive model behavior, and mitigation strategies for advanced AI systems.
- Led the Nova LMArena team, driving Nova models to top-15 LMArena performance by connecting evaluation diagnostics, post-training priorities, and model iteration.
- Founded and led the Nova Trusted AI Challenge, generating 100+ applications and selecting 10 university teams to advance external safety and evaluation research.
- Developed governance-oriented research narratives used with senior leaders and external evaluators, helping Nova models earn recognition from independent safety assessments [1].

Applied Science Manager December 2018 – September 2022
Amazon Alexa / Responsible AI Cambridge, MA

- Founded and led a Responsible AI research team for Alexa, building foundational capabilities for private, fair, and trustworthy customer experiences across platform and product workflows.
- Partnered with product, engineering, and senior executive stakeholders to operationalize responsible AI goals and report outcomes to Amazon Senior Vice Presidents.
- Helped embed responsible AI mechanisms into large-scale conversational AI systems used by broad consumer populations.

Applied Scientist August 2016 – November 2018
Amazon Alexa / Large-Scale NLU Cambridge, MA

- Served as technical lead on decoupling training and testing of large-scale ML systems, improving evaluation rigor for production NLU development.
- Delivered business impact that led to organizational restructuring and influenced production model development practices.

EDUCATION

PhD, Electrical Engineering, University of Southern California 2010 – 2016
Bachelor of Technology, Electrical Engineering, Indian Institute of Technology, Kharagpur 2006 – 2010

EXTERNAL LEADERSHIP & SERVICE

-
- Organizer of the [TrustNLP](#) workshop series and workshops on trustworthy speech processing and LLM unlearning [1], [2]; reviewer for ACL, ICLR, ICASSP, and related top-tier venues.
 - Point of contact for multiple [Amazon Academic Research Awards](#) collaborations, collectively representing over \$20M in funded research.
 - Keynote speaker at venues including [IEEE SLT](#) and workshops on frontier AI advancement and social impact [1], [2].

SELECTED PUBLICATIONS

Full list available at my [Google Scholar page](#).

- [A Threshold Exceedance Framework for CBRN Uplift Evaluation in Frontier Language Models](#), arXiv preprint, 2026. First-author, hands-on experimentation.
- [Establishing Best Practices for Building Rigorous Agentic Benchmarks](#), NeurIPS 2025. Winner of the AgentX LLM Agents MOOC Competition, 2025.
- [FLIRT: Feedback Loop In-context Red Teaming](#), EMNLP 2024. Covered in the red teaming community [1], [2]; #3 paper of the day on Hugging Face.
- [“I’m fully who I am”: Towards centering transgender and non-binary voices to measure biases in open language generation](#), ACM FAccT 2023.
- [BOLD: Dataset and Metrics for Measuring Biases in Open-Ended Language Generation](#), ACM FAccT 2021. (700+ citations)

OTHER RECOGNITIONS

- Ming Hsieh Department of Electrical Engineering best PhD pitch award, 2015.
- Two-time winner of Interspeech computational paralinguistics challenge, 2013 and 2015.
- Phi Beta Kappa alumni in Southern California scholarship, 2015.
- Ming Hsieh Department of Electrical Engineering humanitarian award, 2014.
- Provost PhD fellowship, 2010–2014.